

HW4: Occupation Dataset

Introduction:

Special thanks to: <https://github.com/guipsamora> for sharing his datasets, materials, and questions.

- <https://github.com/justmarkham> for sharing the dataset and materials.

```
In [1]: ### Import the necessary libraries
import pandas as pd
import matplotlib.pyplot as plt
```

```
In [2]: ### Please note that the following data set is from the US in the 20th century.
### Gender proportions for occupations and ages are thus reflective of that era.
### Import the dataset from this address.
### https://raw.githubusercontent.com/justmarkham/DAT8/master/data/u.user
### Assign it to a variable called users and use the 'user_id' as index
users = pd.read_csv('https://raw.githubusercontent.com/justmarkham/DAT8/master/data/u.user',
                    sep='|', index_col='user_id')
```

```
In [3]: # Problem 1. Print the data frame. By default,
# Pandas prints the first few rows and the last few rows.
# This exercise has already been done for you
users
```

```
Out[3]:
```

	age	gender	occupation	zip_code
user_id				
1	24	M	technician	85711
2	53	F	other	94043
3	23	M	writer	32067
4	24	M	technician	43537
5	33	F	other	15213
...
939	26	F	student	33319
940	32	M	administrator	02215
941	20	M	student	97229
942	48	F	librarian	78209
943	22	M	student	77841

943 rows × 4 columns

```
In [4]: # Problem 2. How many observations and columns are in the data?
# Run the necessary command that will output the answer.
print(users.shape[0])
print(users.shape[1])
```

```
943
4
```

```
In [5]: # Problem 3. How many different occupations there are in this dataset?
users.occupation.nunique()
```

```
Out[5]: 21
```

```
In [6]: # Problem 4. What is the most frequent occupation?
users.occupation.mode()
```

```
Out[6]: 0    student
Name: occupation, dtype: object
```

```
In [7]: # Problem 5. Discover what is the mean age per occupation.
# Sort the results and find the 3 occupations with the lowest mean age and the 3 with the highest
mean_age_per_occupation = users.groupby('occupation')['age'].mean().sort_values()

lowest = mean_age_per_occupation.head(3)
highest = mean_age_per_occupation.tail(3)

print(lowest)
print(highest)
```

```

occupation
student      22.081633
none         26.555556
entertainment 29.222222
Name: age, dtype: float64
occupation
educator     42.010526
doctor       43.571429
retired      63.071429
Name: age, dtype: float64

```

```

In [8]: # Problem 6. Find the proportion of males by occupation and sort it from the most to the least
prop_male = users.groupby('occupation')['gender'].apply(lambda x: (x == 'M').mean()).sort_values(ascending = False)
print(prop_male)

```

```

occupation
doctor      1.000000
engineer    0.970149
technician  0.962963
retired     0.928571
programmer  0.909091
executive   0.906250
scientist   0.903226
entertainment 0.888889
lawyer      0.833333
salesman    0.750000
educator    0.726316
student     0.693878
other       0.657143
marketing   0.615385
writer      0.577778
none        0.555556
administrator 0.544304
artist      0.535714
librarian   0.431373
healthcare  0.312500
homemaker   0.142857
Name: gender, dtype: float64

```

```

In [9]: # Problem 7. For each occupation, calculate the minimum and maximum ages
# See groupby and agg() to perform multiple aggregate functions at once

print(users.groupby('occupation')['age'].agg(['min', 'max']))

```

```

      min  max
occupation
administrator  21   70
artist         19   48
doctor         28   64
educator       23   63
engineer       22   70
entertainment  15   50
executive      22   69
healthcare     22   62
homemaker      20   50
lawyer         21   53
librarian      23   69
marketing      24   55
none           11   55
other          13   64
programmer     20   63
retired        51   73
salesman       18   66
scientist      23   55
student        7   42
technician     21   55
writer         18   60

```

```

In [10]: # Problem 8. For each combination of occupation and gender, calculate the mean age.
# Arrange the results in a table so each row is an occupation, and you have a
# column of the average male age and another column with the average female age.
# Sort the resulting table by Female mean age from least to greatest

mean_age = users.groupby(['occupation', 'gender'])['age'].mean().unstack().sort_values(by = 'F')
print(mean_age)

```

gender	F	M
occupation		
student	20.750000	22.669118
salesman	27.000000	38.555556
scientist	28.333333	36.321429
engineer	29.500000	36.600000
artist	30.307692	32.333333
entertainment	31.000000	29.000000
programmer	32.166667	33.216667
homemaker	34.166667	23.000000
other	35.472222	34.028986
none	36.500000	18.600000
marketing	37.200000	37.875000
writer	37.631579	35.346154
technician	38.000000	32.961538
educator	39.115385	43.101449
lawyer	39.500000	36.200000
healthcare	39.818182	45.400000
librarian	40.000000	40.000000
administrator	40.638889	37.162791
executive	44.000000	38.172414
retired	70.000000	62.538462
doctor	NaN	43.571429

```
In [11]: # Problem 9. For each occupation find the count of women and men
# Arrange the results in a table so each row is an occupation, similar to above
count_gender_by_occ = users.groupby('occupation')['gender'].value_counts().unstack()
print(count_gender_by_occ)
```


gender	F	M
occupation		
administrator	36.0	43.0
artist	13.0	15.0
doctor	NaN	7.0
educator	26.0	69.0
engineer	2.0	65.0
entertainment	2.0	16.0
executive	3.0	29.0
healthcare	11.0	5.0
homemaker	6.0	1.0
lawyer	2.0	10.0
librarian	29.0	22.0
marketing	10.0	16.0
none	4.0	5.0
other	36.0	69.0
programmer	6.0	60.0
retired	1.0	13.0
salesman	3.0	9.0
scientist	3.0	28.0
student	60.0	136.0
technician	1.0	26.0
writer	19.0	26.0

```
In [12]: # Problem 10. Turn the counts above into proportions. e.g administrator 0.455696 0.544304
# Arrange results in increasing order of proportion men
sum_gender_by_occ = count_gender_by_occ.sum(axis=1)

prop_gender_per_occ = count_gender_by_occ.div(sum_gender_by_occ, axis=0)
print(prop_gender_per_occ)
```

gender	F	M
occupation		
administrator	0.455696	0.544304
artist	0.464286	0.535714
doctor	NaN	1.000000
educator	0.273684	0.726316
engineer	0.029851	0.970149
entertainment	0.111111	0.888889
executive	0.093750	0.906250
healthcare	0.687500	0.312500
homemaker	0.857143	0.142857
lawyer	0.166667	0.833333
librarian	0.568627	0.431373
marketing	0.384615	0.615385
none	0.444444	0.555556
other	0.342857	0.657143
programmer	0.090909	0.909091
retired	0.071429	0.928571
salesman	0.250000	0.750000
scientist	0.096774	0.903226
student	0.306122	0.693878
technician	0.037037	0.962963
writer	0.422222	0.577778

```
In [13]: # Create a stacked barchart showing the results above
prop_gender_per_occ.plot(kind='bar', stacked = True)
plt.show()
```



```
In [14]: # Extract the first digit of each zip code
# and create a new column called 'region' that maps the
# first digit of the zip to new values using this dictionary:
d = {'0': 'New England',
     '1': 'Mid-Atlantic',
     '2': 'Central East Coast',
     '3': 'The South',
     '4': 'Midwest',
     '5': 'Northern Great Plains',
     '6': 'Central Great Plains',
     '7': 'Southern Central',
     '8': 'Mountain Desert',
     '9': 'West Coast'}

# print the first 5 rows of the result
# postal codes that begin with a letter are actually Canadian but are missing the last digit. These rows can be
```

```
In [15]: users['region'] = users['zip_code'].astype(str).str[0].map(d)
print(users.head(5))
```

	age	gender	occupation	zip_code	region
user_id					
1	24	M	technician	85711	Mountain Desert
2	53	F	other	94043	West Coast
3	23	M	writer	32067	The South
4	24	M	technician	43537	Midwest
5	33	F	other	15213	Mid-Atlantic

```
In [16]: # for the occupation 'retired', find the mean age of each region
retirees = users[users['occupation'] == 'retired']

mean_age_per_region_for_retirees = retirees.groupby('region')['age'].mean()
print(mean_age_per_region_for_retirees)
```

region	age
Central East Coast	60.0
Central Great Plains	59.5
Mid-Atlantic	60.0
Midwest	69.0
New England	65.0
Northern Great Plains	61.0
The South	73.0
West Coast	60.5

Name: age, dtype: float64

In []: